

AI System Inventory Procedure

Classification: STAFF Owner: AI Governance Lead Review date: 2026-12-15 Page 1 of 4

Purpose and scope

CAMT Learning Services uses this fictional document to records approved systems, owners, purposes, and review dates. This document is staff-only operational guidance and must be retrieved only for an authorised staff session. The accountable owner is the AI Governance Lead. The document is reviewed at a planned interval and changes are recorded so that staff and learners can identify the current guidance. Its main topics are inventory, governance, ai. The guidance is written for a teaching demonstration and contains no real student, staff, vendor, or institutional records. The scope is deliberately narrow: it explains the expected action, the audience, and the boundary between published information and staff operational material. A reader should be able to distinguish a general explanation from a decision that requires an accountable human owner. When a request is outside scope, the appropriate response is to name the limitation and route the person to the correct support path rather than inventing an answer.

A practical record for this topic names the decision being made, the information used, the person accountable for the next step, and the review date. The record should be understandable to its intended audience without requiring access to unrelated systems. For the teaching corpus, examples use fictional service labels and aggregate patterns only. A retrieval assistant should quote the relevant passage, preserve its page locator, and state when the record does not support a requested conclusion. This keeps a convenient answer from becoming an unsupported policy decision.

Evidence for AI System Inventory Procedure should remain connected to its stated purpose: records approved systems, owners, purposes, and review dates. A useful answer identifies the relevant rule, the limits of that rule, and the responsible owner rather than merely repeating a familiar keyword. For example, an access question should distinguish a person who may read a published explanation from a person who is authorised to use operational guidance. A review question should identify what evidence is available and what must be referred to an accountable human role. This distinction is important when a language model produces fluent text that appears persuasive but cannot independently verify authority, freshness, or permission.

A fictional classroom scenario asks a user to find information related to inventory, governance, ai. The service first identifies the effective role, then searches only the documents available to that role, and then constructs a bounded evidence bundle. If the requested action is unsupported, conflicts with the document classification, or contains unnecessary sensitive information, the service stops safely and records a minimal reason code. The scenario demonstrates that retrieval quality, access control, output validation, and auditability are separate responsibilities that must work together.

Teaching corpus note: This document is entirely fictional and is included only for secure RAG exercises.

AI System Inventory Procedure

Classification: STAFF Owner: AI Governance Lead Review date: 2026-12-15 Page 2 of 4

Classification and access

This document is staff-only operational guidance and must be retrieved only for an authorised staff session. Access labels are governance metadata, not a visual decoration. A secure retrieval system applies the label before ranking documents and before sending any excerpt to a language model. The user interface may show the effective access level, but the server determines it from an authenticated session. A correct citation cannot repair an access failure after protected text has already been disclosed.

The classification must travel with the document, its chunks, its vector records, and every source link. A public label permits learner-facing discovery, whereas a staff label represents an operational boundary. The application enforces this distinction during retrieval rather than relying on a display control after results have been returned. If a question crosses the boundary, the system should avoid confirming protected document names, identifiers, or wording. It should give a short, useful explanation of the available public scope instead.

Evidence for AI System Inventory Procedure should remain connected to its stated purpose: records approved systems, owners, purposes, and review dates. A useful answer identifies the relevant rule, the limits of that rule, and the responsible owner rather than merely repeating a familiar keyword. For example, an access question should distinguish a person who may read a published explanation from a person who is authorised to use operational guidance. A review question should identify what evidence is available and what must be referred to an accountable human role. This distinction is important when a language model produces fluent text that appears persuasive but cannot independently verify authority, freshness, or permission.

A fictional classroom scenario asks a user to find information related to inventory, governance, ai. The service first identifies the effective role, then searches only the documents available to that role, and then constructs a bounded evidence bundle. If the requested action is unsupported, conflicts with the document classification, or contains unnecessary sensitive information, the service stops safely and records a minimal reason code. The scenario demonstrates that retrieval quality, access control, output validation, and auditability are separate responsibilities that must work together.

Teaching corpus note: This document is entirely fictional and is included only for secure RAG exercises.

AI System Inventory Procedure

Classification: STAFF Owner: AI Governance Lead Review date: 2026-12-15 Page 3 of 4

Operating practice

For this topic, the working practice is to record the purpose of a request, use only the minimum necessary information, and preserve page-level evidence for important answers. Staff should confirm that a source is current, identify the owner when an exception is needed, and avoid placing personal information or confidential operational details in prompts. If evidence is missing or conflicting, the safe response is to request clarification or refer the matter for human review.

Before a workflow is used, the responsible team checks whether the question is in scope, whether the source is current, and whether the evidence is sufficient for the requested action. The team records a concise reason when it declines, escalates, or asks for clarification. This supports review without retaining unnecessary personal information. In this demonstration, an answer is considered grounded only when its citation resolves to a retrieved document that the current role is allowed to access.

Evidence for AI System Inventory Procedure should remain connected to its stated purpose: records approved systems, owners, purposes, and review dates. A useful answer identifies the relevant rule, the limits of that rule, and the responsible owner rather than merely repeating a familiar keyword. For example, an access question should distinguish a person who may read a published explanation from a person who is authorised to use operational guidance. A review question should identify what evidence is available and what must be referred to an accountable human role. This distinction is important when a language model produces fluent text that appears persuasive but cannot independently verify authority, freshness, or permission.

A fictional classroom scenario asks a user to find information related to inventory, governance, ai. The service first identifies the effective role, then searches only the documents available to that role, and then constructs a bounded evidence bundle. If the requested action is unsupported, conflicts with the document classification, or contains unnecessary sensitive information, the service stops safely and records a minimal reason code. The scenario demonstrates that retrieval quality, access control, output validation, and auditability are separate responsibilities that must work together.

Teaching corpus note: This document is entirely fictional and is included only for secure RAG exercises.

AI System Inventory Procedure

Classification: STAFF Owner: AI Governance Lead Review date: 2026-12-15 Page 4 of 4

Review and escalation

The AI Governance Lead reviews this guidance with relevant teaching, privacy, and risk roles. A concern about inaccurate retrieval, unsafe model output, unexpected access, or retained sensitive data is recorded as an incident and assessed using the appropriate procedure. The classroom exercise uses synthetic examples to show that monitoring, audit evidence, and clear responsibilities are part of governance. They are not a substitute for production legal, privacy, or security review.

Review activity includes checking that the document remains relevant, that its classification has not changed, and that known risks have an owner. A concern about access, model behaviour, data handling, or misleading output should be recorded with a reason code and a minimal request fingerprint. It should not require copying raw sensitive text into an audit system. The teaching exercise therefore treats logging as evidence of a decision, not as a hidden archive of all conversations.

Evidence for AI System Inventory Procedure should remain connected to its stated purpose: records approved systems, owners, purposes, and review dates. A useful answer identifies the relevant rule, the limits of that rule, and the responsible owner rather than merely repeating a familiar keyword. For example, an access question should distinguish a person who may read a published explanation from a person who is authorised to use operational guidance. A review question should identify what evidence is available and what must be referred to an accountable human role. This distinction is important when a language model produces fluent text that appears persuasive but cannot independently verify authority, freshness, or permission.

A fictional classroom scenario asks a user to find information related to inventory, governance, ai. The service first identifies the effective role, then searches only the documents available to that role, and then constructs a bounded evidence bundle. If the requested action is unsupported, conflicts with the document classification, or contains unnecessary sensitive information, the service stops safely and records a minimal reason code. The scenario demonstrates that retrieval quality, access control, output validation, and auditability are separate responsibilities that must work together.

Teaching corpus note: This document is entirely fictional and is included only for secure RAG exercises.